
▼ **AIM** : To apply different models on dataset to get value of mean absolute percentage error less than 5% .

DATASET USED : Automobiles

Data Source (Link): <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>

Objective: To get mean absolute percentage error less than 5% by using different machine and deep learning models.

Description: Collected data from UCI Machine Learning repositories. It is a regression problem containing 205 instances and 26 attributes with both numerical and categorical values. This project contains two major parts. First is, cleaning and preprocessing of data. Use of inbuilt functions (like describe and .isnull) to check if any null value is present in data. Conversion of datatype of attributes using .info() .As we have given many attributes; we need to check how much they are correlated to each other for better understanding of data, which can be done using heatmap. We might have lots of missing values, no need to drop all of them; we can also replace null values using aggregates like mean, median, mode, count etc. Standardising also plays main role in preprocessing of data. Use of scatter plot, box and whiskers plot has been done to analyse data more precisely. To know more about data, we used groupby function to play around with data to get useful information. Second is, use of various machine learning and deep learning models to train data and to choose the one giving more accuracy.

▼ Step_1 : Importing libraries and Dataset

```
import pandas as pd
import numpy as np
```

```
# importing data
```

```
#Note :Because the data does not include headers, we can add an argument headers = None inside the read_csv() method, so that pandas will not automatically set the first row as a header
```

```
path = r"https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data"
am_df = pd.read_csv(path, header=None)
```


```
am_df
```

	0	1	2	3	4	5	6	7	8	9	...	16	17	18	19	20	21	22	23	24	25	
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21	27	13495	
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21	27	16500	
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154	5000	19	26	16500	
3	2	164		audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102	5500	24	30	13950
4	2	164		audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115	5500	18	22	17450
...	
200	-1	95		volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114	5400	23	28	16845
201	-1	95		volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	160	5300	19	25	19045
202	-1	95		volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	134	5500	18	23	21485
203	-1	95		volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	23.0	106	4800	26	27	22470

NOTE : Point to Remember : pandas automatically set the header by an integer. To better describe our data we can introduce a header(we create a list "headers" that include all co

headers = ["symboling","normalized-losses","make","fuel-type","aspiration", "num-of-doors","body-style","drive-wheels","engine-location","wheel-base", "length","width","height","cur
print("headers\n", headers)

headers
['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-of-doors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length', 'width', 'height



am_df.columns = headers

am_df

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
2	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	
3	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	

▼ Step_2 : Preprocessing

```

4          2          164      audi      gas      std      four      sedan      4wd      front      99.4      ...      136      mpfi      3.19

```

```

#checking type of data
type(am_df)

```

```

pandas.core.frame.DataFrame

```

```

#checking the datatype for ? using index location
type(am_df.iloc[0,1])

```

```

str

```

```

205 rows x 26 columns

```

```

## NOTE : Relacing all '?' with NaN type
am_df=am_df.replace(['?'],np.nan)

```

```

#The info() method prints information about the DataFrame.

```

```

#The information contains the number of columns, column labels, column data types, memory usage, range index, and the number of cells in each column (non-null values).

```

```

am_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null   int64
1   normalized-losses      164 non-null   object
2   make                   205 non-null   object
3   fuel-type              205 non-null   object
4   aspiration              205 non-null   object
5   num-of-doors           203 non-null   object
6   body-style             205 non-null   object
7   drive-wheels           205 non-null   object
8   engine-location        205 non-null   object
9   wheel-base            205 non-null   float64
10  length                 205 non-null   float64
11  width                  205 non-null   float64
12  height                 205 non-null   float64

```

```

13  curb-weight      205 non-null  int64
14  engine-type      205 non-null  object
15  num-of-cylinders  205 non-null  object
16  engine-size      205 non-null  int64
17  fuel-system      205 non-null  object
18  bore             201 non-null  object
19  stroke            201 non-null  object
20  compression-ratio 205 non-null  float64
21  horsepower       203 non-null  object
22  peak-rpm         203 non-null  object
23  city-mpg         205 non-null  int64
24  highway-mpg      205 non-null  int64
25  price            201 non-null  object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

#converting the object datatype to numeric datatype using the information we get from above code

```
am_df[['price', 'peak-rpm', 'horsepower', 'stroke', 'bore', 'normalized-losses']] = am_df[['price', 'peak-rpm', 'horsepower', 'stroke', 'bore', 'normalized-losses']].apply(pd.to_num
```

am_df.info() #performing this code again to check if dtype for variables has changed or not

```

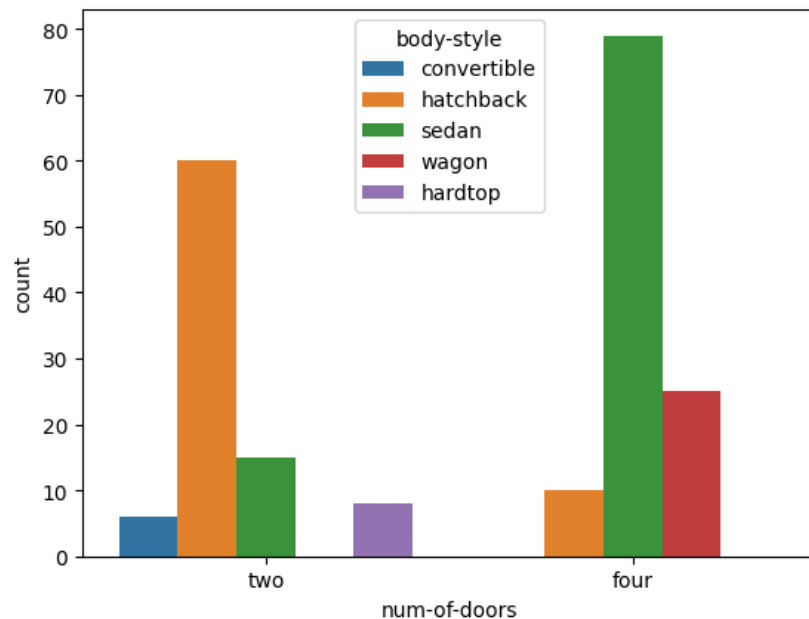
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   symboling            205 non-null    int64
1   normalized-losses    164 non-null    float64
2   make                 205 non-null    object
3   fuel-type            205 non-null    object
4   aspiration           205 non-null    object
5   num-of-doors         203 non-null    object
6   body-style           205 non-null    object
7   drive-wheels         205 non-null    object
8   engine-location      205 non-null    object
9   wheel-base          205 non-null    float64
10  length               205 non-null    float64
11  width                205 non-null    float64
12  height               205 non-null    float64
13  curb-weight          205 non-null    int64
14  engine-type          205 non-null    object
15  num-of-cylinders     205 non-null    object
16  engine-size          205 non-null    int64
17  fuel-system          205 non-null    object
18  bore                 201 non-null    float64
19  stroke               201 non-null    float64
20  compression-ratio    205 non-null    float64
21  horsepower           203 non-null    float64
22  peak-rpm             203 non-null    float64
23  city-mpg             205 non-null    int64
24  highway-mpg          205 non-null    int64
25  price                201 non-null    float64

```

```
dtypes: float64(11), int64(5), object(10)
memory usage: 41.8+ KB
```

```
import seaborn as sns
sns.countplot(data=am_df, x='num-of-doors', hue='body-style')
```

<Axes: xlabel='num-of-doors', ylabel='count'>



```
am_df['num-of-doors'].mode()
```

```
0    four
Name: num-of-doors, dtype: object
```

```
#Filling misssing values(im columns having num values)
```

```
for j in am_df.columns:
    if am_df[j].dtype != 'object':

        am_df[j] = am_df[j].fillna(am_df[j].mean())
```

```
am_df['num-of-doors'].value_counts()
```

```
four    114
two      89
Name: num-of-doors, dtype: int64
```

```
am_df.loc[:, 'num-of-doors'].fillna('four', inplace=True) # (filling categorical variable: 'num-of-doors' with value 'four' in place of null values)
```

```
am_df.isnull().any() # Checking if there is any null value left
```

```
symboling      False
normalized-losses  False
make           False
fuel-type      False
aspiration     False
num-of-doors   False
body-style     False
drive-wheels   False
engine-location False
wheel-base    False
length        False
width         False
height        False
curb-weight    False
engine-type    False
num-of-cylinders False
engine-size    False
fuel-system    False
bore          False
stroke        False
compression-ratio False
horsepower     False
peak-rpm       False
city-mpg       False
highway-mpg    False
price         False
dtype: bool
```

▼ Step_2.1: Separating X(input variables) and Y(target/output variable)(i.e 'price')

```
X = am_df.iloc[:, 0:-1]
Y = am_df.iloc[:, -1]
```

```
X_list_num = [j for j in X.columns if X[j].dtype != 'object']    ##List of numeric columns headers # will be needing this list later on in this project
X_list_cat = [j for j in X.columns if X[j].dtype == 'object']    ##List of category column headers
```

▼ Step_3 : Standardising X

```
#Note : y can not b standardised because this is Linear Regression problem
```

```
# Standardising numerical variables using MinMaxScaler which scales all values in between a range of 0-1
```

```

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

#fit_transform() - calculates mean and standard dev of each column - (x-mean)/std
X_numerical_scaled = scaler.fit_transform(X[X_list_num].to_numpy())    ## .to_numpy converts dataframe to numpy array

```

X_numerical_scaled

```

array([[1.          , 0.29842932, 0.05830904, ..., 0.34693878, 0.22222222,
        0.28947368],
       [1.          , 0.29842932, 0.05830904, ..., 0.34693878, 0.22222222,
        0.28947368],
       [0.6          , 0.29842932, 0.2303207 , ..., 0.34693878, 0.16666667,
        0.26315789],
       ...,
       [0.2          , 0.15706806, 0.65597668, ..., 0.55102041, 0.13888889,
        0.18421053],
       [0.2          , 0.15706806, 0.65597668, ..., 0.26530612, 0.36111111,
        0.28947368],
       [0.2          , 0.15706806, 0.65597668, ..., 0.51020408, 0.16666667,
        0.23684211]])

```

X_numerical_scaled.shape, X[X_list_num].shape

```
((205, 15), (205, 15))
```

#converting numpy array back to dataframe

```

X_numerical_scaled = pd.DataFrame(X_numerical_scaled, columns = X_list_num)
X_numerical_scaled.head(2)

```

	symboling	normalized- losses	wheel- base	length	width	height	curb- weight	engine- size	bore	stroke	compression- ratio	horsepower
0	1.0	0.298429	0.058309	0.413433	0.316667	0.083333	0.411171	0.260377	0.664286	0.290476	0.125	0.2625

#creating dummies for one hot encoding categorical variables

```
X_categorical_ohe = pd.get_dummies(data = X[X_list_cat], columns = X_list_cat)
```

X['fuel-system'].unique() #checking no. of classes present in column specified in code

```

array(['mpfi', '2bbl', 'mfi', '1bbl', 'spfi', '4bbl', 'idi', 'spdi'],
      dtype=object)

```

X_categorical_ohe.head()

	make_alfa-romero	make_audi	make_bmw	make_chevrolet	make_dodge	make_honda	make_isuzu	make_jaguar	make_mazda	make_mercedes-benz
0	1	0	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	0	0
2	1	0	0	0	0	0	0	0	0	0
3	0	1	0	0	0	0	0	0	0	0
4	0	1	0	0	0	0	0	0	0	0

5 rows × 60 columns



```
X = pd.concat([X_categorical_ohe, X_numerical_scaled,], axis = 1)
```

X #input variables are now finally preprocessed

	make_alfa-romero	make_audi	make_bmw	make_chevrolet	make_dodge	make_honda	make_isuzu	make_jaguar	make_mazda	make_mercedes ben
0	1	0	0	0	0	0	0	0	0	
1	1	0	0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	0	0	
3	0	1	0	0	0	0	0	0	0	
4	0	1	0	0	0	0	0	0	0	
...
200	0	0	0	0	0	0	0	0	0	
201	0	0	0	0	0	0	0	0	0	
202	0	0	0	0	0	0	0	0	0	
203	0	0	0	0	0	0	0	0	0	
204	0	0	0	0	0	0	0	0	0	

205 rows × 75 columns



```
type(X)
```

```
pandas.core.frame.DataFrame
```

```
type(Y)
```



```
pandas.core.series.Series
```

▼ Step_3.1 : converting to numpy array

```
import numpy as np
```

```
X = np.array(X)
```

```
Y = np.array(Y)
```

```
X.shape
```

```
(205, 75)
```

```
Y.shape
```

```
(205,)
```

▼ Step_4 : Splitting into test and train data

```
from sklearn.model_selection import train_test_split
```

```
Xtrain, Xtest, ytrain, ytest = train_test_split(X, Y , test_size=0.2, random_state=42, shuffle=True)
```

```
Xtrain.shape
```

```
(164, 75)
```

```
Xtest.shape
```

```
(41, 75)
```

▼ Step_5 : Building ANN model

```
#Step_5.1 : Importing models
```

```
from keras.models import Sequential
```

```
from keras.layers import Dense
```

```
am_df_ANN = Sequential()
```

```

#Step_5.2 : Adding layers

'''
Note: sigmoid activation is for - binary [2 class classification]
      softmax activation is for - multi [2+ class classification]
'''

'''
Note: output layer: number of units = number of class/category present in target variable (for multi [2+ class classification])
      : number of units = 1 (for binary [2 class classification])
'''

#hidden layer
am_df_ANN.add(Dense(units=750, activation = 'relu'))
#hidden layer
am_df_ANN.add(Dense(units=300, activation = 'relu'))
#output layer
am_df_ANN.add(Dense(units=1, activation = 'relu'))

#Step_5.3 : Exponential decay, Compile, best mmodel, early stopping

from tensorflow.keras.optimizers.schedules import ExponentialDecay
from keras.optimizers import Adam
'''
ExponentialDecay - with iterations reduce the learning rate
'''

initial_learning_rate = 0.001
lr=ExponentialDecay(initial_learning_rate,
                    decay_steps=100000,
                    decay_rate=0.96,
                    staircase=True)

am_df_ANN.compile(loss='mean_squared_error',
                  metrics='mean_absolute_percentage_error',
                  optimizer='adam'
                  )

'''
Note: categorical_crossentropy - loss for multiclass classification
      binary_crossentropy - loss for binary classification
      mean_squared_error - loss for linear regression problems
'''

from keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau

'''
EarlyStopping- is used when metric doesn't improve for certain epochs we stop training

EarlyStopping(
    monitor="val_accuracy", - metric to monitor
    patience=10, - number of epochs we monitor metric

```

```
verbose=1, - message will be printed)
...

es = EarlyStopping(monitor="val_accuracy", patience=10, verbose=1)
mc = ModelCheckpoint(filepath='bestmodel.h5', monitor='val_accuracy', mode='max', verbose=1, save_best_only=True)
```

▼ Step_6 : Fit

```
history = am_df_ANN.fit(Xtrain, ytrain, epochs=150)
```

```
Epoch 116/150
6/6 [=====] - 0s 12ms/step - loss: 6488474.0000 - mean_absolute_percentage_error: 14.5695
Epoch 117/150
6/6 [=====] - 0s 12ms/step - loss: 6425537.0000 - mean_absolute_percentage_error: 14.5188
Epoch 118/150
6/6 [=====] - 0s 12ms/step - loss: 6364487.5000 - mean_absolute_percentage_error: 14.4803
Epoch 119/150
6/6 [=====] - 0s 14ms/step - loss: 6302905.0000 - mean_absolute_percentage_error: 14.4599
Epoch 120/150
6/6 [=====] - 0s 15ms/step - loss: 6234747.5000 - mean_absolute_percentage_error: 14.3643
Epoch 121/150
6/6 [=====] - 0s 12ms/step - loss: 6195698.5000 - mean_absolute_percentage_error: 14.2191
Epoch 122/150
6/6 [=====] - 0s 20ms/step - loss: 6186924.5000 - mean_absolute_percentage_error: 14.1227
Epoch 123/150
6/6 [=====] - 0s 15ms/step - loss: 6061656.0000 - mean_absolute_percentage_error: 14.0217
Epoch 124/150
```

#Step_6.1 : ANN architecture

```
am_df_ANN.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 750)	57000
dense_1 (Dense)	(None, 300)	225300
dense_2 (Dense)	(None, 1)	301
Total params: 282,601		
Trainable params: 282,601		
Non-trainable params: 0		


#Step_6.2 : Plotting loss and mean absolute percentage error

```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['loss'])
plt.show()
```



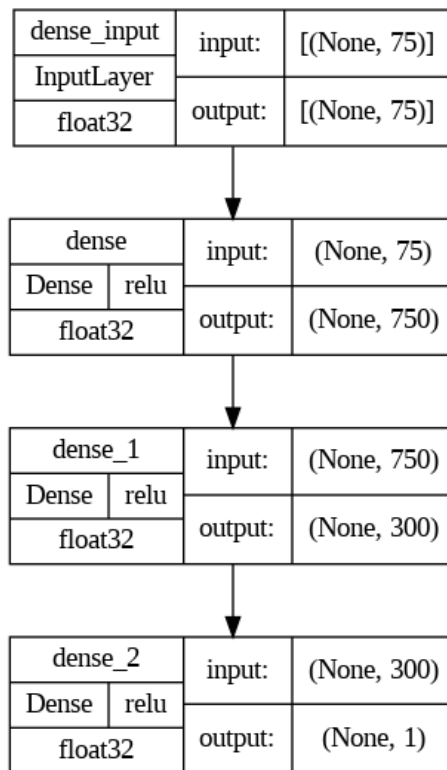
```
plt.plot(history.history['mean_absolute_percentage_error'])  
plt.show()
```



▼ Step_7 : Model plotting

```
#conda install -c anaconda pydot
#conda install -c anaconda graphviz
```

```
from tensorflow.keras.utils import plot_model
plot_model(am_df_ANN, show_shapes=True, show_dtype=True, show_layer_activations=True, show_layer_names=True)
```



▼ Step_8 : Prediction of ypred using X

```
ypred = am_df_ANN.predict(Xtest)
```

```
2/2 [=====] - 0s 11ms/step
```

```
ypred
```

```
ypred.shape
```

```
(41, 1)
```

```
am_df_ANN.evaluate(Xtest, ytest)
```

```
2/2 [=====] - 0s 8ms/step - loss: 12898391.0000 - mean_absolute_percentage_error: 20.8886  
[12898391.0, 20.888553619384766]
```

We can see that ANN model is not precisely built and giving us 20% error but our aim is to get error less than 5%. Therefore we will now be applying Decision Trees model to get better accuracy/less error.

▼ Decision Trees

```
from sklearn.tree import DecisionTreeRegressor  
from sklearn.model_selection import GridSearchCV
```

Note : Grid Search uses a different combination of all the specified hyperparameters and their values and calculates the performance for each combination and selects the best value for the hyperparameters. This makes the processing time-consuming and expensive based on the number of hyperparameters involved.

#visit this link to know more about hyperparameters of decision tree : <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeRegressor.html#sklearn.tree.DecisionTreeRegressor>

```
#hyper-parameters of decision trees  
'''
```

```
criterion = ['squared_error', 'friedman_mse', 'absolute_error', 'poisson'] #calculating the error of prediction  
splitter = ['best','random']  
max_depth = [1, 5, 10, None]  
min_samples_split = [2, 3, 5, 10, 15]  
min_samples_leaf = [1, 2, 3, 5, 10]  
min_weight_fraction_leaf = [0, 2, 5, 10]  
max_features = [None, 'auto', 'sqrt', 'log2']  
max_leaf_nodes = [None, 2, 3, 4, 5]  
min_impurity_decrease = [0,1,2,3]  
ccp_alpha = [0,1,2,3]
```

```
#this will take alot of time ...so we are now reducing some values  
'''
```

```
criterion = ['friedman_mse', 'absolute_error'] #calculating the error of prediction  
splitter = ['best']  
max_depth = [1, 5, 10, None]  
min_samples_split = [2, 3, 5]  
min_samples_leaf = [1, 2, 3]  
min_weight_fraction_leaf = [0, 2, 5]  
max_features = [None, 'auto', 'sqrt', 'log2']  
max_leaf_nodes = [None, 2, 5]  
min_impurity_decrease = [0,1,3]  
ccp_alpha = [0,1,3]
```

```
param_dict = dict(criterion = criterion,
                  splitter= splitter,
                  max_depth = max_depth,
                  min_samples_split = min_samples_split,
                  min_samples_leaf = min_samples_leaf,
                  min_weight_fraction_leaf = min_weight_fraction_leaf,
                  max_features = max_features,
                  max_leaf_nodes = max_leaf_nodes,
                  min_impurity_decrease = min_impurity_decrease,
                  ccp_alpha =ccp_alpha)
```

```
obj = DecisionTreeRegressor()
```

```
grid = GridSearchCV(estimator=obj, param_grid=param_dict, scoring = 'neg_root_mean_squared_error', verbose=1)
```

```
grid_result = grid.fit(Xtrain, ytrain)
```

```
finalmodel = grid_result.best_estimator_#it will show us the best value for each hyperparameter which is being chosen by gridsearch
```

```
finalmodel
```

```
DecisionTreeRegressor
DecisionTreeRegressor(ccp_alpha=1, criterion='absolute_error',
                      max_features='sqrt', min_impurity_decrease=3,
                      min_weight_fraction_leaf=0)
```

```
ytestpred = finalmodel.predict(Xtest)
```

```
from sklearn.metrics import mean_absolute_percentage_error
```

```
mean_absolute_percentage_error(ytest, ytestpred)
```

```
0.1519439298166267
```

Therefore; Accuracy = (100-0.15)% = 99.85%

▼ Random Forests

```
from sklearn.ensemble import RandomForestRegressor
```



```
obj1 = RandomForestRegressor()
```

```
n_estimators=[2, 5, 10]#how many decision trees we want ?
```

```
criterion = ['friedman_mse'] #calculating the error of prediction
max_depth = [1, 5, 10, None]
min_samples_split = [2, 5]
min_samples_leaf = [1, 3]
min_weight_fraction_leaf = [0, 2, 5]
max_features = [None,'sqrt', 'log2']
max_leaf_nodes = [None, 2, 5]
min_impurity_decrease = [0,1,3]
ccp_alpha = [0,1,3]
```

```
param_dict = dict(n_estimators=n_estimators,
                  criterion = criterion,
                  max_depth = max_depth,
                  min_samples_split = min_samples_split,
                  min_samples_leaf = min_samples_leaf,
                  min_weight_fraction_leaf = min_weight_fraction_leaf,
                  max_features = max_features,
                  max_leaf_nodes = max_leaf_nodes,
                  min_impurity_decrease = min_impurity_decrease,
                  ccp_alpha =ccp_alpha)
```

```
grid1 = GridSearchCV(estimator=obj1, param_grid=param_dict, scoring = 'neg_root_mean_squared_error', verbose=1)
```

```
grid_result1 = grid1.fit(Xtrain, ytrain)
```

```
Fitting 5 folds for each of 11664 candidates, totalling 58320 fits
/usr/local/lib/python3.9/dist-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
38880 fits failed out of a total of 58320.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

```
Below are more details about the failures:
```

```
-----
19440 fits failed with the following error:
```

```
Traceback (most recent call last):
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/ensemble/_forest.py", line 340, in fit
    self._validate_params()
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/base.py", line 600, in _validate_params
    validate_parameter_constraints(
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/utils/_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
```

```
sklearn.utils._param_validation.InvalidParameterError: The 'min_weight_fraction_leaf' parameter of RandomForestRegressor must be a float in the range [0.0, 0.5]. Got 2 instead
```

19440 fits failed with the following error:

Traceback (most recent call last):

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/ensemble/_forest.py", line 340, in fit
    self._validate_params()
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/base.py", line 600, in _validate_params
    validate_parameter_constraints(
```

```
File "/usr/local/lib/python3.9/dist-packages/sklearn/utils/_param_validation.py", line 97, in validate_parameter_constraints
    raise InvalidParameterError(
```

```
sklearn.utils._param_validation.InvalidParameterError: The 'min_weight_fraction_leaf' parameter of RandomForestRegressor must be a float in the range [0.0, 0.5]. Got 5 instead
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/model_selection/_search.py:952: UserWarning: One or more of the test scores are non-finite: [-5611.77267003 -4685.95464249 -4840
nan nan]
    warnings.warn(
```

```
finalmodel1 = grid_result1.best_estimator_
```

```
finalmodel1
```

```
RandomForestRegressor
RandomForestRegressor(ccp_alpha=1, criterion='friedman_mse', max_depth=10,
                      max_features='log2', min_impurity_decrease=1,
                      min_samples_split=5, min_weight_fraction_leaf=0,
                      n_estimators=10)
```

```
ytestpred1 = finalmodel1.predict(Xtest)
```

```
from sklearn.metrics import mean_absolute_percentage_error
```

```
mean_absolute_percentage_error(ytest, ytestpred1)
```

```
0.13359206261878148
```

Therefore; Accuracy = (100-0.13)% = 99.87%

CONCLUSION : From all three model applied above i.e Artificial Neural Network(ANN), Decision Trees, Random Forest; we can say that Random Forest with accuracy of 99.85 is best model of all.

